

Fine Tuning LLMs

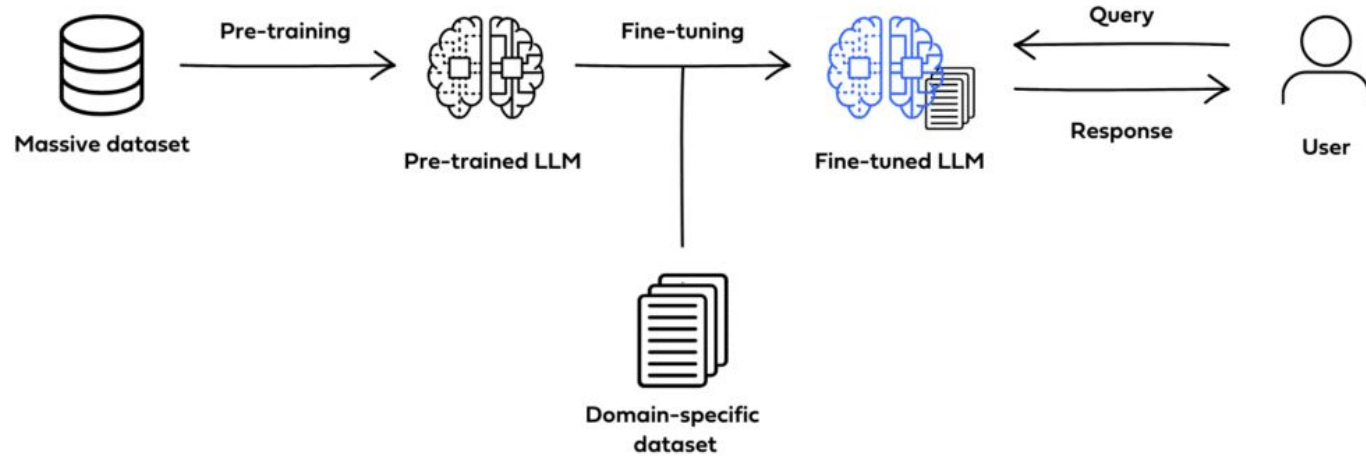
Spurthi Setty

Why should you finetune



- More Consistent Outputs
- Customize models for specific use cases
- Reduces Hallucinations
- Eliminates need of training a model from scratch

Fine Tuning Architecture




Fine Tuning Methods

- Self-Supervised Learning
 - Unlabeled Data → more scalable
 - Model will mimic the style of the text
- Supervised Learning
 - Training data consists of inputs and outputs
 - Can train for a specific task via instruction fine tuning
 - Ex. Translation, summarization, question-answering
- Reinforcement Learning with Human Feedback
 - Develop a reward model with model outputs and human ratings
 - Proximal Policy Optimization (PPO):



Input	Output



Dataset Creation and Processing

- Each LLM requires a specific format for the training data
 - Ex. Stanford Alpaca format for LLAMA-2
 - GPT requires a specific data format specified in their docs
 - Usually in a json or jsonl format
- High quality data essential for better performance
- Tokenize the formatted data
 - Converts text to numbers
- Train-Test split

```
sample = ""
```

```
### Instructions: {instruction}
```

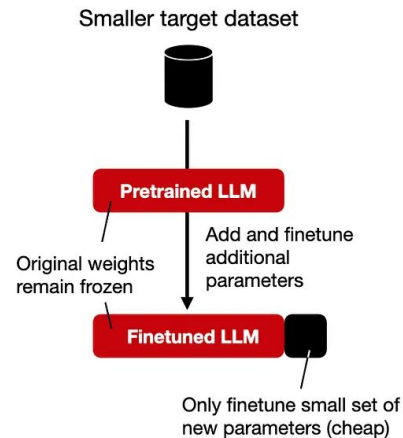

```
### Input: {Input}
```

```
### Response:
```

```
""
```


Methods for Parameter Fine Tuning

- Full Fine Tuning
 - Adjusts all parameters of the LLM using task-specific data.
 - Computationally expensive
- Transfer Learning
 - Freeze all parameters except for the head of the neural network
 - Only finetune the layers that translate to the output layer
- Parameter Efficient Fine Tuning (PEFT)
 - Freeze all the weights of the base LLM
 - Augment the model with additional parameters and finetune those
 - Less computationally expensive



Low Rank Adaptation (LoRA) for PEFT

- Reduces number of trainable parameters
- Identifies the most crucial parameters for the task at hand and finetunes those
 - approximate the weight matrices of these important layers using low-rank matrices.
- During fine-tuning, only the parameters in the low-rank matrices are updated
- Less chance of overfitting since only a few parameters are updated
- Reduces the computational and memory requirements needed for fine-tuning



Integrating multiple techniques to improve performance

- Prompt Engineering and RAG can be used in conjunction with Fine Tuning
- Use other methods to create optimal training data for Fine Tuning
 - Prompt engineering to optimize system prompts for a specific task
 - Ex. “You are a financial expert.. ”
 - Use prompt engineering to guide outputs to be a certain way
 - Can include context as a part of the input data that is retrieved via RAG
- Use prompt engineering and RAG on a fine tuned model the same way one does on a base model
 - Specialized knowledge from RAG
 - Control and specificity from prompt engineering
- Combination of techniques needed for highly sophisticated, accurate, and efficient AI systems

Thank you!